

EDUCATION

Massachusetts Institute of Technology

Ph.D. in Computer Science

Advisor: Antonio Torralba

Cambridge, MA

2020–Current

Massachusetts Institute of Technology

M.S. in Computer Science

Cambridge, MA

2017–2020

– GPA: 5.00/5.00

– Thesis: “Automatic Relational Data Augmentation for Machine Learning”

– Hewlett Packard Fellowship

Rutgers University

B.S. in Computer Science

New Brunswick, NJ

2013–2016

– GPA : 4.00/4.00

– Summa Cum Laude graduation honor

– Academic Excellence Scholarship (2015)

– Academic Excellence Award (2014)

– Winner of Database challenge project (2014)

EXPERIENCE

Microsoft Research

Software Engineering Intern in Research

Redmond, WA

Summer 2020

– Worked in the DMX group on AutoML research project FLAML.

– Improved and added various system components in the system to increase the speed and accuracy of prediction.

– Worked with regression, SVM, random forest, XGBoost, Catboost, neural models, and other predictive models, methods, and algorithms.

– Improved sampling techniques.

– Incorporated ensemble models in the pipeline.

– Tuned various hyper-parameters for improved performance.

– Tested and verified correctness of computational pipelines.

– Identified and fixed bugs.

Adobe Research

Software Engineering Intern in Research

San Jose, CA

Summer 2018

– Worked on Question-Answering (NLP) system for Creative Cloud.

– Designed and implemented neural models for Question-Answering on text.

– Used language modeling as unsupervised pre-training.

– Built a demo system with model integration.

Rutgers University

Software Engineer and Research Fellow

New Brunswick, NJ

Jan 2016-Jan 2017

– Integrated various APIs and services that are currently available in the form of different mobile applications.

- Combined the chain of APIs and Deeplinks to unify multiple services under one interactive chat interface.
- Developed a web crawler to mine app descriptions and analyze the content.
- Designed stratified sampling techniques for Apache Spark, including how to maximize precision for strata under fixed sample size, and how to sample the data under user-provided error bound for estimating aggregation functions such as sum and mean.
- Benchmarked sampling in Apache Spark framework versus Hadoop framework.

CareCloud

Software Engineering Intern

Miami, FL

Summer 2015

- Optimized queries and procedures for medical reports.
- Migrated and extended pivot table for multidimensional database (OLAP cube).
- Developed server health-check service.
- Gathered information from third-party software providers.

Luxoft

QA Engineer

Kyiv, Ukraine

Sept 2012-Aug 2013

- Customer: UBS Bank
- Worked in FICC department with CDF derivatives.
- Automated routines to verify correctness of daily reports to clearing house.
- Tested various financial trading platforms.

SELECTED PROJECTS

Dissecting Backpropagation

Interpreting how the learning dynamics of a deep network evolves over the course of the training phase.

- Developed a new *network dissection* method to visualize and quantify the emergence of semantics in representation gradients during training.
- Studied overfitting and adversarial robustness through the lens of interpretability.
- Studied emergent backpropagation neuron features with VQGAN+CLIP.

Generator-based Adversarial Attacks

Creating generator-based adversarial attacks for vision models that are transferable between different architectures and datasets in black-box fashion.

- Mentoring (UROP): Sauhaarda (Raunak) Chowdhuri (MIT EECS).
- Developed a generator-based adversarial attacks with novel learning techniques that incorporate deep model ensembles and produce cross-domain, black-box, and universal attacks in milliseconds.
- Introduced diversity induced objective approach and improved generator-based attacks in cross-domain and black-box fashion by 10-15%.

Automatic Relational Data Augmentation

Finding new features relevant to the user's predictive task with minimal "human-in-the-loop" involvement.

- Designed and implemented an end-to-end system that takes as input a data repository, and outputs an augmented data set such that training a predictive model on this augmented dataset results in improved prediction accuracy.
- Designed novel feature selection randomization based algorithm that outperformed other feature selection algorithms in highly noisy setup.
- Outperformed multiple AutoML systems performance in both accuracy and time.
- Experimented with targeted structural similarity based attacks to decrease attack's visual perceptibility.

TEACHING

- **Teaching Assistant** at MIT Fall 2019
Advances in Computer Vision (6.819/6.869)
- **Teaching Assistant** at Get Smarter Spring 2019
Human Computer Interaction

PROGRAMMING

- Languages
Python, Java, C, JavaScript, SQL, HTML/CSS
- Libraries and Tools
PyTorch, Scikit-learn, Pandas, SciPy, JUnit

TALKS

1. “**Weighted Maximum Independent Set of Geometric Objects in Turnstile Streams**”. APPROX 2020
2. “**Robust and Sample Optimal Algorithms for PSD Low-Rank Approximation**”. FOCS 2020
3. “**ARDA: Automatic Relational Data Augmentation for Machine Learning**”. VLDB 2020

PUBLICATIONS IN SYSTEMS AND MACHINE LEARNING

1. **Learning Program Representations for Food Images and Cooking Recipes.**
Dim P. Papadopoulos, Enrique Mora, Nadiia Chepurko, Kuan Wei Huang, Ferda Ofli, Antonio Torralba
CVPR 2022.
2. **ARDA: Automatic Relational Data Augmentation for Machine Learning.**
Nadiia Chepurko, Ryan Marcus, Emmanuel Zggragen, Raul Castro Fernandez, Tim Kraska, David Karger
VLDB 2020.

PUBLICATIONS IN ALGORITHMS

1. **Near-Optimal Algorithms for Linear Algebra in the Current Matrix Multiplication Time.**
Nadiia Chepurko*, Kenneth L. Clarkson*, Praneeth Kacham*, Lior Horesh*, David Woodruff*
SODA 2022.
2. **Quantum-Inspired Algorithms from Randomized Numerical Linear Algebra.**
Nadiia Chepurko*, Kenneth L. Clarkson*, Lior Horesh*, Honghao Lin*, David Woodruff*
ICML 2022.
3. **Robust and Sample Optimal Algorithms for PSD Low-Rank Approximation.**
Nadiia Chepurko*, Ainesh Bakshi*, David Woodruff*
FOCS 2020.
4. **Testing Positive Semi-Definiteness via Random Submatrices.**
Nadiia Chepurko*, Ainesh Bakshi*, Rajesh Jayaram*
FOCS 2020.
5. **Weighted Maximum Independent Set of Geometric Objects in Turnstile Streams.**
Nadiia Chepurko*, Ainesh Bakshi*, David Woodruff*
APPROX 2020.

*Equal contribution